

Chapter 2

DBMS Architecture

Data Models – Categories of data models, Schemas, Instances, and Database state. DBMS Architecture and Data Independence – The Three schema architecture, Data independence, DBMS Languages and Interfaces , Classifications of Database Management Systems.

Data Models

A **data model** is a collection of concepts that can be used to describe the structure of a database. By structure of a database means the data types, relationships, and constraints that apply to the data. Most data models also include a set of basic operations for specifying retrievals and updates on the database.

Data Abstraction is the suppression of details of data organization and storage, and the highlighting of the essential features for an improved understanding of data(simply means general features). In order to achieve the data abstraction data model is used.

Categories of data models

- Hight level or Conceptual data model
- Low level or Physical data model
- Representational or Implementational data model

High-level or conceptual data models provide concepts that are close to the way many users perceive data. Conceptual data models use concepts such as **entities, attributes, and relationships**. An **entity** represents a real-world object or concept, such as an employee or a project from the miniworld that is described in the database. An **attribute** represents some property of interest that further describes an entity, such as the employee's name or salary. The **relationship** among two or more entities represents an association among the entities, for example, a works-on relationship between an employee and a project.

Low-level or physical data models describe how data is stored as files in the computer by representing information such as record formats, record orderings, and access paths. An access path is a search structure that makes the search for particular database records efficient, such as indexing or hashing. An index is an example of an access path that allows direct access to data using an index term or a keyword. It is similar to the index at the end of this text, except that it may be organized in a linear, hierarchical (tree-structured), or some other fashion.

Representational Data Model this type of data model is used to represent only the logical part of the database and does not represent the physical structure of the database. The representational data model allows us to focus primarily, on the design part of the database. A popular representational model is a **Relational model**. The relational Model consists of

Relational Algebra and **Relational Calculus**. In the Relational Model, we basically use tables to represent our data and the relationships between them. It is a theoretical concept whose practical implementation is done in Physical Data Model.

The advantage of using a Representational data model is to provide a foundation to form the base for the Physical model

Another class of data models is known as **self-describing data models**. The data storage in systems based on these models combines the description of the data with the data values themselves. In traditional DBMSs, the description (schema) is separated from the data. These models include XML as well as many of the key-value stores and NOSQL systems that were recently created for managing big data.

Schemas, Instances, and Database state

In data model, it is important to distinguish between the **description** of the database and the database itself. The description of a database is called the **database schema**, which is specified during database design and is not expected to change frequently. Most data models have certain conventions for displaying schemas as diagrams. A displayed schema is called a **schema diagram**. Each object in the schema—such as STUDENT or COURSE—a **schema construct**

STUDENT			
Name	Student_number	Class	Major
COURSE			
Course_name	Course_number	Credit_hours	Department
PREREQUISITE			
Course_number	Prerequisite_number		
SECTION			
Section_identifier	Course_number	Semester	Year
GRADE_REPORT			
Student_number	Section_identifier	Grade	

Figure 1:Schema Diagram for Student database

STUDENT

Name	Student_number	Class	Major
Smith	17	1	CS
Brown	8	2	CS

COURSE

Course_name	Course_number	Credit_hours	Department
Intro to Computer Science	CS1310	4	CS
Data Structures	CS3320	4	CS
Discrete Mathematics	MATH2410	3	MATH
Database	CS3380	3	CS

SECTION

Section_identifier	Course_number	Semester	Year	Instructor
85	MATH2410	Fall	07	King
92	CS1310	Fall	07	Anderson
102	CS3320	Spring	08	Knuth
112	MATH2410	Fall	08	Chang
119	CS1310	Fall	08	Anderson
135	CS3380	Fall	08	Stone

GRADE_REPORT

Student_number	Section_identifier	Grade
17	112	B
17	119	C
8	85	A
8	92	A
8	102	B
8	135	A

PREREQUISITE

Course_number	Prerequisite_number
CS3380	CS3320
CS3380	MATH2410
CS3320	CS1310

Figure 2 Student Database

A schema diagram displays only some aspects of a schema, such as the names of record types and data items, and some types of constraints. Other aspects are not specified in the schema diagram; for example, Figure 1 shows neither the data type of each data item nor the relationships among the various files. Many types of constraints are not represented in schema diagrams.

The actual data in a database may change quite frequently. For example, the database shown in Figure 2 changes every time we add a new student or enter a new grade. The data in the database at a particular moment in time is called a **database state or snapshot**. It is also called the **current set of occurrences or instances** in the database. In a given database state, each schema construct has its own current set of instances; for example, the STUDENT construct will contain the set of individual student entities (records) as its instances. Many database states can be constructed to correspond to a particular database schema. Every time we insert or delete a record or change the value of a data item in a record, we change one state of the database into another state.

The distinction between **database schema** and **database state** is very important. When we define a new database, we specify its **database schema** only to the DBMS. At this point, the corresponding database state is the **empty state with no data**. We get the initial state of the database when the database is **first populated or loaded with the initial data**. From then on, every time an update operation is applied to the database, we get another database state. At any point in time, the database has a current state. The DBMS is partly responsible for ensuring that every state of the database is a valid state—that is, a state that satisfies the structure and constraints specified in the schema. Hence, specifying a correct schema to the DBMS is extremely important and the schema must be designed with utmost care. The DBMS stores the descriptions of the schema constructs and constraints called the **meta-data**—in the DBMS catalog so that DBMS software can refer to the schema whenever it needs to. The schema is sometimes called the **intension**, and a database state is called an **extension** of the schema.

Although, as mentioned earlier, the schema is not supposed to change frequently, it is not uncommon that changes occasionally need to be applied to the schema as the application requirements change. For example, we may decide that another data item needs to be stored for each record in a file, such as adding the Date_of_birth to the STUDENT schema in Figure 1. This is known as **schema evolution**. Most modern DBMSs include some operations for schema evolution that can be applied while the database is operational.

DBMS Architecture and Data Independence

An architecture for database systems, called the three-schema architecture, that proposed to help achieve and visualize some of the main characteristics of database system. They are

- Use of a catalog to store the database description (schema) so as to make it self-describing
- Insulation of programs and data
- Support of multiple user views

The Three-Schema Architecture

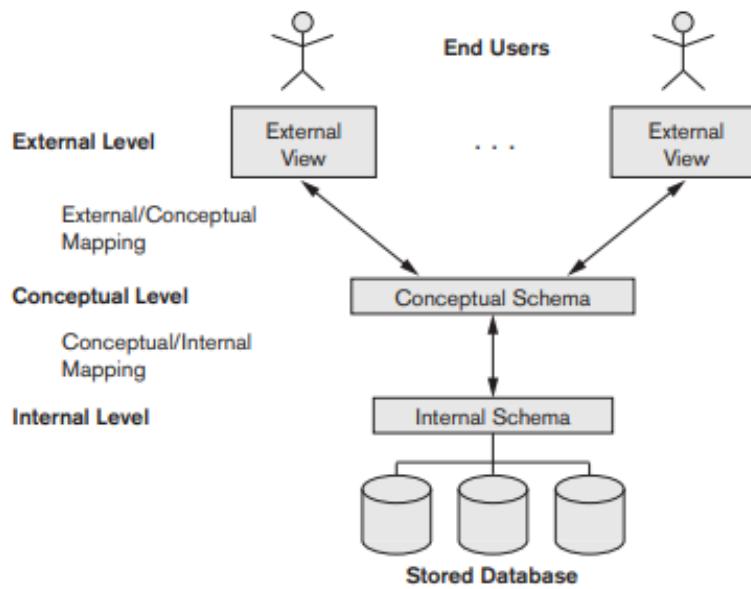


Figure 3 Schema Architecture

The goal of the three-schema architecture, illustrated in Figure 3, is to separate the user applications from the physical database. In this architecture, schemas can be defined at the following three levels:

1. The internal level has an internal schema, which describes the physical storage structure of the database. The internal schema uses a physical data model and describes the complete details of data storage and access paths for the database
2. The conceptual level has a conceptual schema, which describes the structure of the whole database for a community of users. The conceptual schema hides the details of physical storage structures and concentrates on describing entities, data types, relationships, user operations, and constraints. Usually, a representational data model is used to describe the conceptual schema when a database system is implemented. This implementation conceptual schema is often based on a conceptual schema design in a high-level data model.
3. The external or view level includes a number of external schemas or user views. Each external schema describes the part of the database that a particular user group is interested in and hides the rest of the database from that user group. As in the previous level, each external schema is typically implemented using a representational data model, possibly based on an external schema design in a high-level conceptual data model.

The three schemas are only descriptions of data; the actual data is stored at the physical level only. In the three-schema architecture, each user group refers to its own external schema. Hence, the DBMS must transform a request specified on an external schema into a request against the conceptual schema, and then into a request on the internal schema for processing over the stored database. If the request is a database retrieval, the data extracted from the stored database must be reformatted to match the user's external view. The processes of transforming requests and results between levels are called **mappings**. These mappings may be time-consuming, so some DBMSs—especially those that are meant to

support small databases—do not support external views. Even in such systems, however, it is necessary to transform requests between the conceptual and internal levels.

Data Independence

The three-schema architecture can be used to further explain the concept of data independence, which can be defined as the capacity to change the schema at one level of a database system without having to change the schema at the next higher level. We can define two types of data independence:

1. Logical data independence is the capacity to change the conceptual schema without having to change external schemas or application programs. We may change the conceptual schema to expand the database (by adding a record type or data item), to change constraints, or to reduce the database (by removing a record type or data item). In the last case, external schemas that refer only to the remaining data should not be affected. Only the view definition and the mappings need to be changed in a DBMS that supports logical data independence. After the conceptual schema undergoes a logical reorganization, application programs that reference the external schema constructs must work as before. Changes to constraints can be applied to the conceptual schema without affecting the external schemas or application programs

EXAMPLE:

If there is a database of a banking system and we want to add the details of a new customer or we want to update or delete the data of a customer at the logical level data will be changed but it will not affect the Physical level or structure of the database.

These changes can be done at a logical level without affecting the application program or external layer.

- Adding, deleting, or modifying the entity or relationship.
- Merging or breaking the record present in the database.

2 Physical data independence is the capacity to change the internal schema without having to change the conceptual schema. Hence, the external schemas need not be changed as well. Changes to the internal schema may be needed because some physical files were reorganized—for example, by creating additional access structures—to improve the performance of retrieval or update. If the same data as before remains in the database, we should not have to change the conceptual schema. For example, providing an access path to improve retrieval speed of SECTION records (Figure 2) by semester and year should not require a query such as list all sections offered in fall 2008 to be changed, although the query would be executed more efficiently by the DBMS by utilizing the new access path.

Below changes can be done at the physical layer without affecting the conceptual layer -

- Changing the storage devices like SSD, hard disk and magnetic tapes, etc.
- Changing the access technique and modifying indexes.
- Changing the compression techniques or hashing algorithms.

Data independence occurs because when the schema is changed at some level, the schema at the next higher level remains unchanged; only the mapping between the two levels is changed. Hence, application programs referring to the higher-level schema need not be changed.

Physical Data Independence	Logical Data Independence
Physical data independence is used to change the internal schema without requiring a change in the logical schema.	Logical data independence is making sure that if you add any new field or delete any existing field we do not need to change the application program.
Physical data independence is easy to attain in comparison to logical data independence.	It is difficult to attain logical data independence compared to physical data independence.
Physical data independence provides feasibility if we want to shift the database or want to change the file organization structure.	Logical data independence helps us to change the data definition and the structure of the data without having changes in the physical schema.
Physical data independence deals with the internal structure of the schema.	Logical data independence deals with conceptual schema.
Examples of changes in Physical independence are Changing the compression techniques, hashing algorithms, SSD, location of the database, etc.	Examples of changes in logical independence are Adding, deleting, or modifying the entity or relationship.

DBMS Languages

Databases serve an important function to organize and store information on a computer. In programming, developers use database languages to maintain and monitor an electronic

database and its management system. These languages perform a variety of critical tasks that help a database management system function correctly

Database languages, also known as query languages or data query languages, are a classification of programming languages that developers use to define and access databases, which are collections of organized data that users can access electronically. These languages allow users to complete tasks such as controlling access to data, defining and updating data and searching for information within the database management system (DBMS). A DBMS is a piece of technology that interacts with users, applications and the database to record and analyze data while also manipulating the database to offer a way to store, access and retrieve data.

A DBMS provides necessary database languages that allow users to express database updates and queries, which are requests for data. There are different examples of database languages available, including SQL, which is the standard programming language for many databases. Database languages comprise four sublanguages that serve different functions to execute tasks.

Categories of database languages

Here are four types of database languages and their uses:

1. Data definition language (DDL)

Data definition language (DDL) creates the framework of the database by specifying the database schema, which is the structure that represents the organization of data. Its common uses include the creation and alteration of tables, files, indexes and columns within the database. This language also allows users to rename or drop the existing database or its components. Here's a list of DDL statements:

- **CREATE:** Creates a new database or object, such as a table, index or column
- **ALTER:** Changes the structure of the database or object
- **DROP:** Deletes the database or existing objects
- **RENAME:** Renames the database or existing objects

2. Data manipulation language (DML)

Data manipulation language (DML) provides operations that handle user requests, offering a way to access and manipulate the data that users store within a database. Its common functions include inserting, updating and retrieving data from the database. Here's a list of DML statements:

- **INSERT:** Adds new data to the existing database table
- **UPDATE:** Changes or updates values in the table
- **DELETE:** Removes records or rows from the table
- **SELECT:** Retrieves data from the table or multiple tables(**DQL -Data Query Language**)

3. Data control language (DCL)

Data control language (DCL) controls access to the data that users store within a database. Essentially, this language controls the rights and permissions of the database system. It allows users to grant or revoke privileges to the database. Here's a list of DCL statements:

- **GRANT:** Gives a user access to the database
- **REVOKE:** Removes a user's access to the database

4. Transaction control language (TCL)

Transaction control language (TCL) manages the transactions within a database. Transactions group a set of related tasks into a single, executable task. All the tasks must succeed in order for the transaction to work. Here's a list of TCL statements:

- **COMMIT:** Carries out a transaction
- **ROLLBACK:** Restores a transaction if any tasks fail to execute
- **SAVEPOINT:** Sets a point in a transaction to save

DBMS Interfaces

User-friendly interfaces provided by a DBMS may include the following

Menu-based Interfaces for Web Clients or Browsing. These interfaces present the user with lists of options (called menus) that lead the user through the formulation of a request. Menus do away with the need to memorize the specific commands and syntax of a query language; rather, the query is composed step-by-step by picking options from a menu that is displayed by the system. Pull-down menus are a very popular technique in Web-based user interfaces. They are also often used in browsing interfaces, which allow a user to look through the contents of a database in an exploratory and unstructured manner.

Apps for Mobile Devices. These interfaces present mobile users with access to their data. For example, banking, reservations, and insurance companies, among many others, provide apps that allow users to access their data through a mobile phone or mobile device. The apps have built-in programmed interfaces that typically allow users to log in using their account name and password; the apps then provide a limited menu of options for mobile access to the user data, as well as options such as paying bills (for banks) or making reservations (for reservation Web sites).

Forms-based Interfaces. A forms-based interface displays a form to each user. Users can fill out all of the form entries to insert new data, or they can fill out only certain entries, in which case the DBMS will retrieve matching data for the remaining entries. Forms are usually designed and programmed for naive users as interfaces to canned transactions. Many DBMSs have forms specification languages, which are special languages that help programmers specify such forms. SQL*Forms is a form-based language that specifies queries using a form designed in conjunction with the relational database schema. Oracle Forms is a component of

the Oracle product suite that provides an extensive set of features to design and build applications using forms. Some systems have utilities that define a form by letting the end user interactively construct a sample form on the screen.

Graphical User Interfaces. A GUI typically displays a schema to the user in diagrammatic form. The user then can specify a query by manipulating the diagram. In many cases, GUIs utilize both menus and forms.

Keyword-based Database Search. These are somewhat similar to Web search engines, which accept strings of natural language (like English or Spanish) words and match them with documents at specific sites (for local search engines) or Web pages on the Web at large (for engines like Google or Ask). They use predefined indexes on words and use ranking functions to retrieve and present resulting documents in a decreasing degree of match. Such “free form” textual query interfaces are not yet common in structured relational databases, although a research area called keyword-based querying has emerged recently for relational databases.

Speech Input and Output. Limited use of speech as an input query and speech as an answer to a question or result of a request is becoming commonplace. Applications with limited vocabularies, such as inquiries for telephone directory, flight arrival/departure, and credit card account information, are allowing speech for input and output to enable customers to access this information. The speech input is detected using a library of predefined words and used to set up the parameters that are supplied to the queries. For output, a similar conversion from text or numbers into speech takes place.

Interfaces for Parametric Users. Parametric users, such as bank tellers, often have a small set of operations that they must perform repeatedly. For example, a teller is able to use single function keys to invoke routine and repetitive transactions such as account deposits or withdrawals, or balance inquiries. Systems analysts and programmers design and implement a special interface for each known class of naive users. Usually a small set of abbreviated commands is included, with the goal of minimizing the number of keystrokes required for each request.

Interfaces for the DBA. Most database systems contain privileged commands that can be used only by the DBA staff. These include commands for creating accounts, setting system parameters, granting account authorization, changing a schema, and reorganizing the storage structures of a database.

Classifications of Database Management Systems

1. Based on Data Model
2. Based on Number of Users
3. Based on Database Distribution
4. Based on Cost of Database

5. Based on Usage
6. Based on Flow Control

1. Based on Data Model

The data model defines the physical and logical structure of a database which involves the data types, the relationship among the data, constraints applied on the data and even the basic operations specifying retrieval and updation of data in the database. Depending upon how the data is structured, data models are further classified into:

a. Relational Data Model

In the relational data model, we use tables to represent data and the relationship among that data. Each of the tables in the relational data model has a unique name. A table has multiple columns where each column name is unique. A table holds records which have values for each column of the table.

We also refer to the relational database model as a record-based data model as it holds records of fixed-format. The relational database model is the most currently used data model.

b. Entity-Relationship Model

The Entity-Relationship model (E-R data model) represents data using objects and the relationship among these objects. These objects are referred to as entities that represent the real ‘thing’ or ‘object’ in the real world. Each entity in the E-R model is distinguishable from other entities in the model. Like, relational model, the E-R model is also used widely used to design the database.

c. Object-Based Data Model

Nowadays, object-oriented programming such as Java, C++, etc. is widely used to develop most of the software. This motivated the development of an object-based data model. The object-based data model is an extension of the E-R model which also includes the notion of encapsulation, methods. There is also an object-relational data model which is a combination of the object-oriented data model and relational data model.

d. Semistructured Data Model

The semistructured data model is different from what we have studied above. In the semistructured data model, the data items or objects of the same kind might have a different set of attributes. The Extensible Markup Language represents the semistructured data.

Before the above data model was introduced, the data models like hierarchical data model and network data model were used to design a database that is still used in some legacy applications. The hierarchical data model stores the data in the form of records and uses a tree structure to represent these records. The record is arranged in

a tree structure where there is a single parent record for each child record. To overcome the shortcomings of the hierarchical data model network data model was introduced which allow the multiple parent record for a single child record.

2. Classification Based on Number of Users

The database management system can also be classified on the basis of its user. So, a DBMS can either be used by a single user or it can be used by multiple users. The database system that can be used by a single user at a time is referred to **as a single-user system** and the database system that can be used by multiple users at a time is referred to as a **multiple user system**.

3. Based on Database Distribution

Depending on the distribution of the database over numerous sites we can classify the database as:

a. Centralized DBMS

In the centralized DBMS, the entire database is stored in a single computer site. Though the centralized database support multiple users still the DBMS software and the data both are stores on a single computer site.

b. Distributed DBMS

In the distributed DBMS (DDBMS) the database and the DBMS software are distributed over many computer sites. These computer sites are connected via a computer network. The DDBMS is further classified as homogeneous DDBMS and heterogeneous DDBMS.

- **Homogeneous DDBMS:** The homogeneous DDBMS has the same DBMS software at all the distributed sites.
- **Heterogeneous DDBMS:** The heterogeneous DDBMS has different DBMS software for different sites.

4. Based on Cost of Database

Well, it is quite difficult to classify the database on the basis of its cost as nowadays you can have free open source DBMS products such as MySQL and PostgreSQL. Although the personal version of RDBMS can cost up to \$100.

The large systems along with the components that can handle distribution of database, replication of database, parallel processing, mobile capability and so on can be sold in the form of licenses. The site license allows unlimited use while another kind of license limits the number of concurrent licenses. Some systems with single-user versions such as Microsoft Access are sold per copy or included in the configuration of your desktop or laptop.

You may also have to pay millions of dollars for the installation and maintenance of a large database system.

5. Classification Based on Usage

On the basis of the access path that is used to store the files, the database can be classified as general-purpose DBMS and special-purpose DBMS. The **special-purpose DBMS** is the one that is designed for a specific application and it can not be used for another application without performing any major changes we refer to this as **online transaction processing (OLTP)**. The OLTP system supports a large number of transactions concurrently without any delay.

The **general-purpose** DBMS is the one that is designed to meet the need of as many applications as possible.

6. Based on Flow Control

Based upon the flow of control from application to DBMS the database management system is broadly classified into two types active database management system and passive database management system. Let us discuss each in brief.

With the **passive database management system**, the user needs to specify the query to the current state of the database system to retrieve the desired information. It is similar to traditional DBMS where the user or the application program is responsible to initiate the operation. As the application forwards the request to DBMS and waits for DBMS to process the query or requested operation and provide a possible answer.

The requested operation may be used to define or update the schema of the database or retrieve or update the data within the database. The passive database management systems are also referred to as program driven systems.

The **active database management system** on other hand are referred to as data-driven systems or event-driven systems where the control flow between the application and DBMS is based on the occurrence of an event. Inactive DBMS the control flow is two ways i.e. application can call DBMS and even DBMS can call application.

Inactive DBMS the users need to specify what data they require. The DBMS processes the request and if the requested information is currently available then it provides it to the user. If the requested data is not available currently then inactive DBMS the scope of query also includes the future data i.e., the DBMS will monitor the arrival of desired information in future data and will provide it to the relevant user.

